

AIVisib: how the tool measures a brand's visibility in AI answers

For partners, agencies and clients. It describes how the product works, the measurement method, the statistical guarantees and security, without confidential implementation details.

1 • What the tool does

- Measures, every week, whether a brand is recommended by ChatGPT, Gemini, Perplexity and Claude when a customer asks a real question (“which dental clinic in Muscat?”).
- Delivers: visibility score, share of recommendations versus tracked competitors, visibility per engine and per language, average position, sentiment, sources cited by the AIs, executive summary and action plan.
- Native languages: English, French, Modern Standard Arabic (questions, answers, reports and interface). Other languages on request.

2 • Architecture (overview)

- **Website and client application:** static web pages served by a global CDN (HTTPS, security headers, CSP).
- **Measurement engine:** Node.js service on a managed cloud platform, automatic deployment on every update, scheduled jobs for weekly measurements.
- **Database:** managed PostgreSQL with per-account isolation (row-level access rules); raw results are only accessible to the service, never directly to the browser.
- **Payments:** Stripe (subscriptions, customer portal, invoices); no card data passes through or is stored by AIVisib.
- **AI providers:** official APIs of OpenAI, Google, Perplexity and Anthropic, called in a neutral state (no history, no personalisation), with live web search.

3 • Measurement method

- **Questions:** generated from brand, sector, city and languages, spread across intents (best choice, near me, price, service, reputation, situational); questions that are too similar are automatically removed so that only genuinely different questions remain, then each one is checked against reality before they count; the client can edit them.
- **Querying:** each question is asked to each engine several times (passes); calls run in parallel and are retried on transient errors; an engine that could not be measured is never shown as 0%.
- **Answer analysis:** detection of the brand mention and position, other recommended brands, tone and cited sources; three anti-hallucination layers (structured extraction, cross-check, deterministic guards).
- **Frequency:** first measurement on demand (a few minutes), then automatic weekly measurement; each report is compared with the previous one (gained / lost by engine, language and question).

4 • Statistical guarantees

- Per question: **Wilson 95%** confidence interval on the proportion of answers naming the brand.
- Per report: the question is the sampling unit; overall confidence is estimated by **bootstrap resampling** over questions, reflecting the real stability of the score rather than false precision. Every score is shown with its margin ($\pm$).
- **Head vs long-tail**: generic and specific questions are measured separately, because the specific ones are what a small business can win first.
- Methodology audited by a third party (August 2026); all audit recommendations implemented.

5 • Security and privacy

- Account authentication (email + password, or Google); each account only sees its own projects; separate client spaces for agencies with revocable invitations.
- Read-only share links, revocable, no account needed for the reader, white-labelled on Pro and Agency plans.
- Client data is never sold, shared or used to train models; projects or accounts can be deleted at any time.
- Encrypted transport (HTTPS/HSTS), secrets kept out of the code, error logging without personal data, rate limiting on public endpoints.

6 • Deliverables and integrations

- Web dashboard, **vector A4 PDF** export (server-side rendering, right-to-left Arabic), shared page, automatic weekly report email to the end client.
- White label (agency name and logo) on reports, PDFs and shared pages from the Pro plan.
- Public free audit: 3 questions × 2 engines live, no account needed.
- No public API yet (planned); data export on request.

7 • Plans

- Starter \$29/month (1 brand, 25 questions) • Pro \$99/month (3 brands, 100 questions, white label) • Agency \$249/month (10 brands, 400 questions, unlimited seats, client links, automatic emails). All engines and languages included, no paid add-ons; cancel in one click.